

# Containers (with emphasis on Docker)

Jakub Daněš  
Yoso Czech s.r.o.

<mailto:jakub.danek@yoso.fi>

# What are containers?

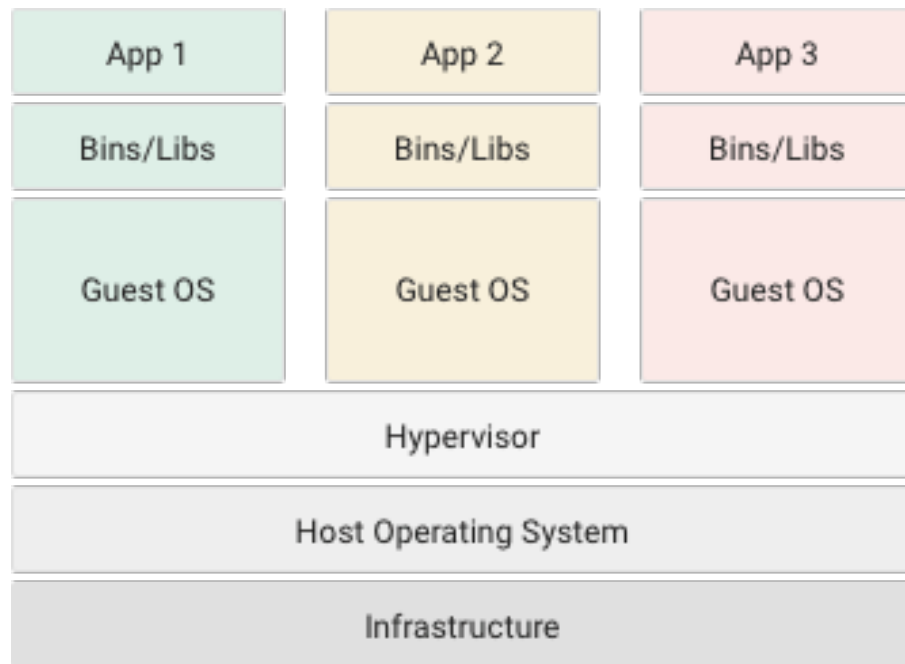
- **Abstraction of applications from the actual runtime environment**
- **Consistent deployment independent on system**
  - Cloud
  - Laptop
  - Own servers

# What are containers?

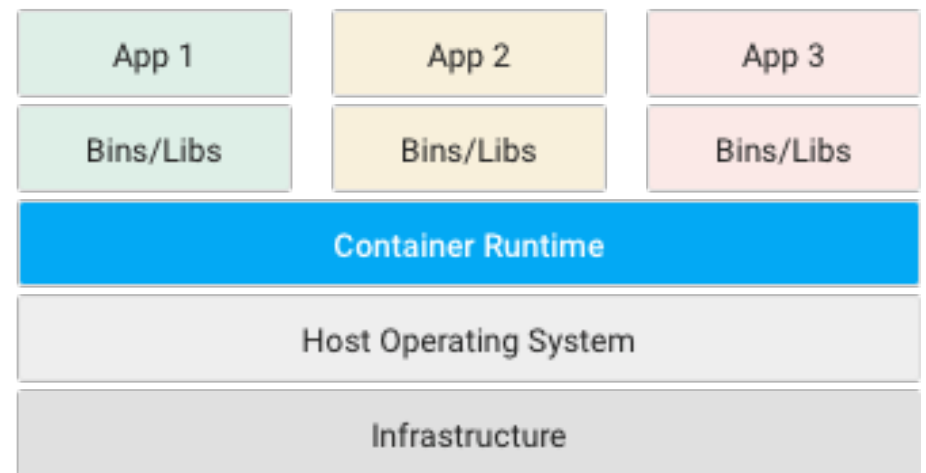
- **Benefits of Consistency**

- Run anywhere
  - Relatively easy to run on many different platforms
- Same environment in dev, test and production
  - → easier transmission from one to another
  - → less bugs introduced by environment differences

# Container vs Virtual Machine



Virtual Machines



Containers

# Container vs Virtual Machine

## Virtual Machines

- Hardware virtualization → one physical server becomes many virtual servers
- Each VM contains full operating system
- Each VM has own (possibly different kernel)
- Each VM has to boot during start (init scripts, systemd...)

# Container vs Virtual Machine

## Containers

- Process-level virtualization → each container runs as process group in host OS
- All containers on single host use the same kernel
- Process startup is much faster than VM boot (generally)
- Enabled by multiple kernel features for process constraint and isolation

# Cgroups

## Control groups

- Linux kernel feature allowing setting resource limits for a process group
- Tree-structure
- Each process belongs to exactly one cgroup
  - All threads spawned by a process belong to the same cgroup as the parent process


# Namespaces

- Means of isolation of processes running on the same host.
- Tree structure
- Processes in one namespace do not see resources of another namespace
- But processes in parent namespace can see resources in child namespaces



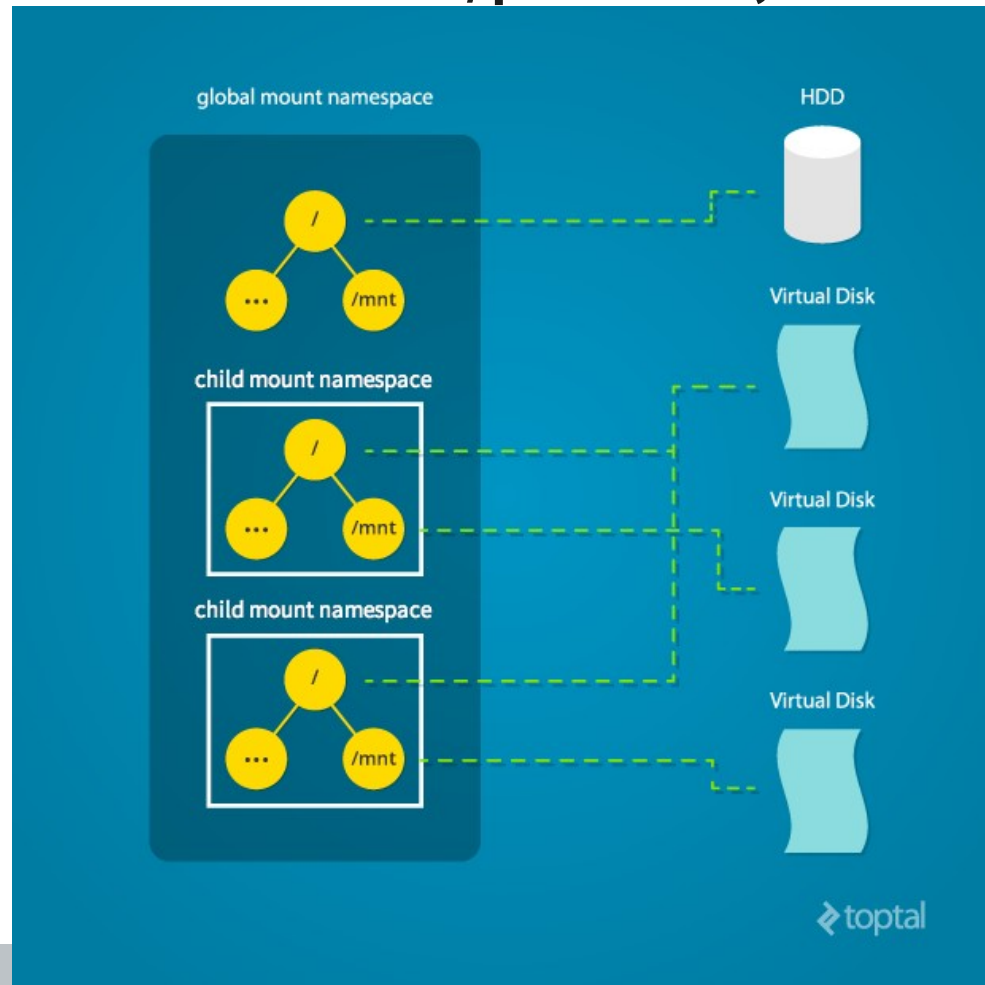
# Namespaces - PID

- Process has different ID for each PID namespace to which it belongs



# Namespaces - Mount

- When isolated process mounts new disk/partition, it affects only its namespace



# Namespaces - types

- cgroups
- user IDs (UID)
- network
- mount points
- interprocess communication (IPC)
- process IDs (PID)
- hostname (UTS)

# And some more

- **Seccomp**

- Preventing processes from making particular system calls

- **Capabilities**

- Permissions to certain actions generally allowed only to privileged users (root)
- Can be enabled per-thread for unprivileged processes

- **Apparmor and SELinux**

- Some more program restricting :)



# More reading

## Linux Man pages:

<http://man7.org/linux/man-pages/man7/capabilities.7.html>

<https://www.kernel.org/doc/Documentation/cgroup-v2.txt>

<http://man7.org/linux/man-pages/man7/namespaces.7.html>

## Under the hood of Docker:

<https://pasztor.at/blog/under-the-hood-of-docker>

## Namespaces:

<https://www.toptal.com/linux/separation-anxiety-isolating-your-system-with-linux-name-spaces>

## Rami Rosen (the page is not pretty, but contents are awesome):

<http://ramirose.wixsite.com/ramirosen>

- **API for creating linux containers**
- **In general, the containers run full OS**
  - Including init, cron, syslog...
- **→ similar result as with VMs, but implemented as containers**
  - Keeps flexibility of VMs
  - Keeps speed and lower overhead of containers

# LXD

- Container manager built with LXC
- Provides common images in remote repository
- Scalability, security, resource management
- Easy to get into (comparing to low-level LXC API)
- <https://linuxcontainers.org/lxd/try-it/>

# Docker

- **Another container engine**
- **Younger than LXC**
- **Different approach: single-application deployment**
  - → docker containers do not run full OS
  - → provide isolated environment for single-application runtime



# Docker

- **First Class Citizens:**

- Image - binary form of the runtime environment
  - identified by id, name and tag (version)
- Container - running environment based on an image
  - identified by unique ID and name


- **Images are stored in a repository (private, public, local)**

- similar mechanism to Maven - if you do not have an image locally, docker will download it for you from remote repositories

# Building Docker Images

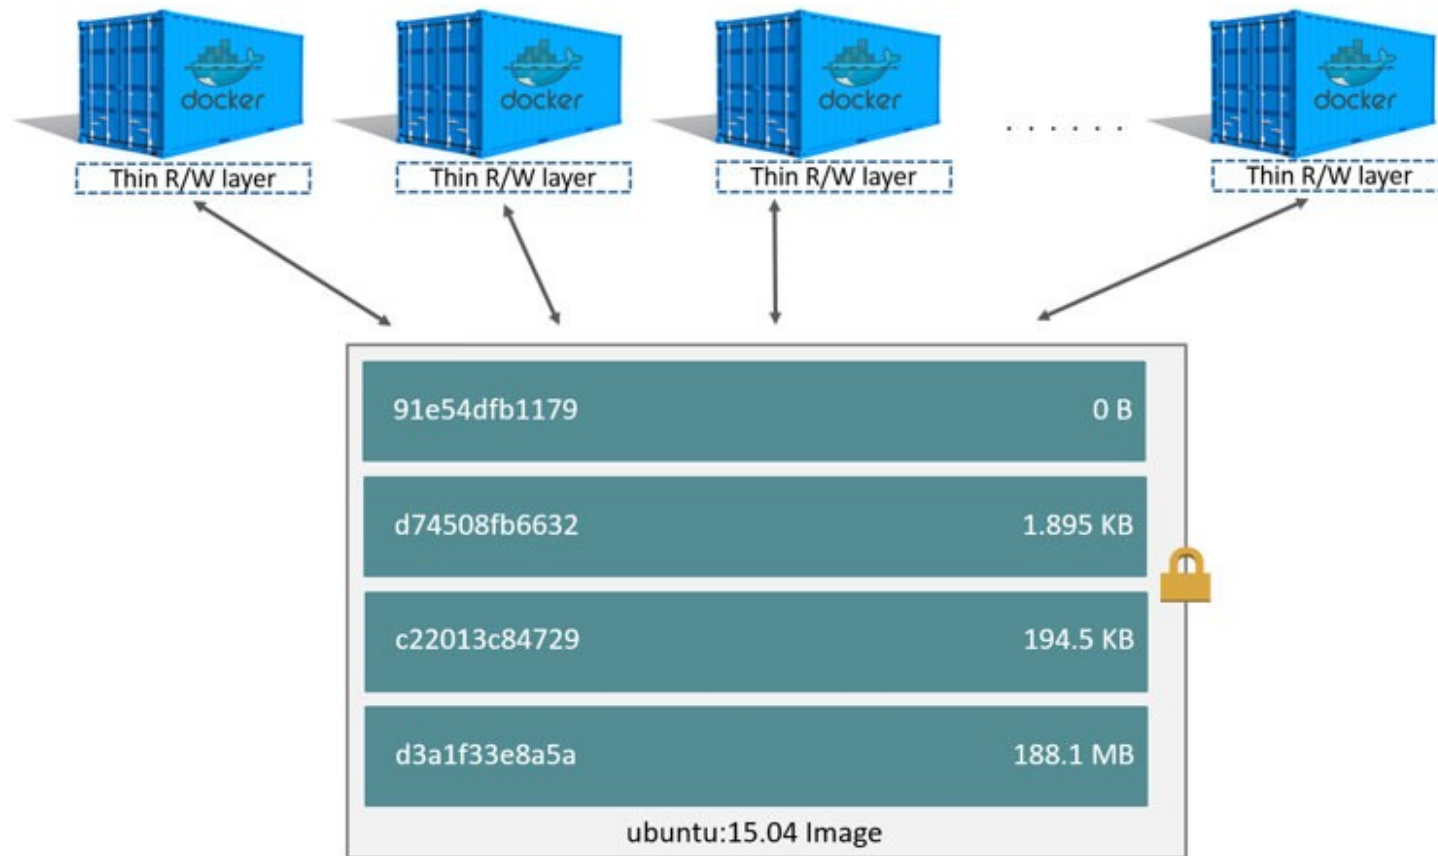
- Images defined by text specification (Dockerfile)
- Images are built of a series of filesystem layers

```
FROM ubuntu:15.04  
COPY . /app  
RUN make /app  
CMD python /app/app.py
```



# Building Docker Images

- Containers have own R/W layer → can share same image



# Docker Containers

- **State of container (changes to the R/W layer) can be committed to create new image**
- **Until then they need to be considered transient**
  - → not good for storing application data
- **Filesystem layers introduce increased complexity to the filesystem implementation**
- **When an update is needed, new image is built and new container started (and old one is killed)**
  - → different philosophy to LXC or VMs where updates are done in runtime

# Docker – Storage Options

- **Bind Mounts**

- older way of storing persistent data
- mount a folder on host system to a folder inside container
- -> you need to know which folder you want to mount
  - it must exist before starting the container

- **Example: C:\\Users\\danekja\\mysql -> /var/lib/mysql**

- When container runs mysql and it stores data, they can be found on host filesystem

# Docker – Storage Options

- **Volumes**

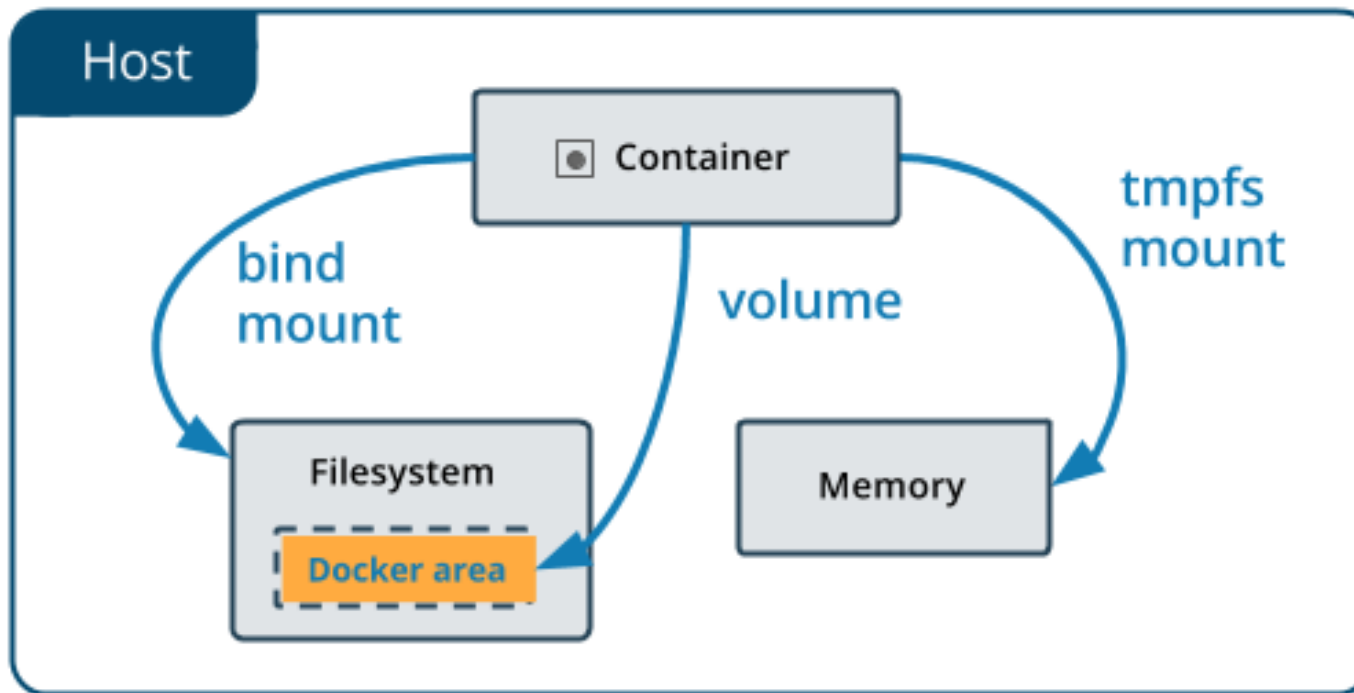
- Volumes are created on host filesystem and mounted to a path inside the container
- Difference to bind mounts: volumes are created by docker in space managed by docker (usually in /var/lib/docker on Linux hosts)
  - on Windows hosts this is a bit complicated because Docker actually runs in Linux VM

# Docker – Storage Options

- **Volumes**

- Volumes can be created independently of containers
  - as most Docker objects are identified by hash ID and a name
- -> Volumes can be shared among multiple containers
- Volumes can be remote

# Docker – Storage Options





# Docker - networking

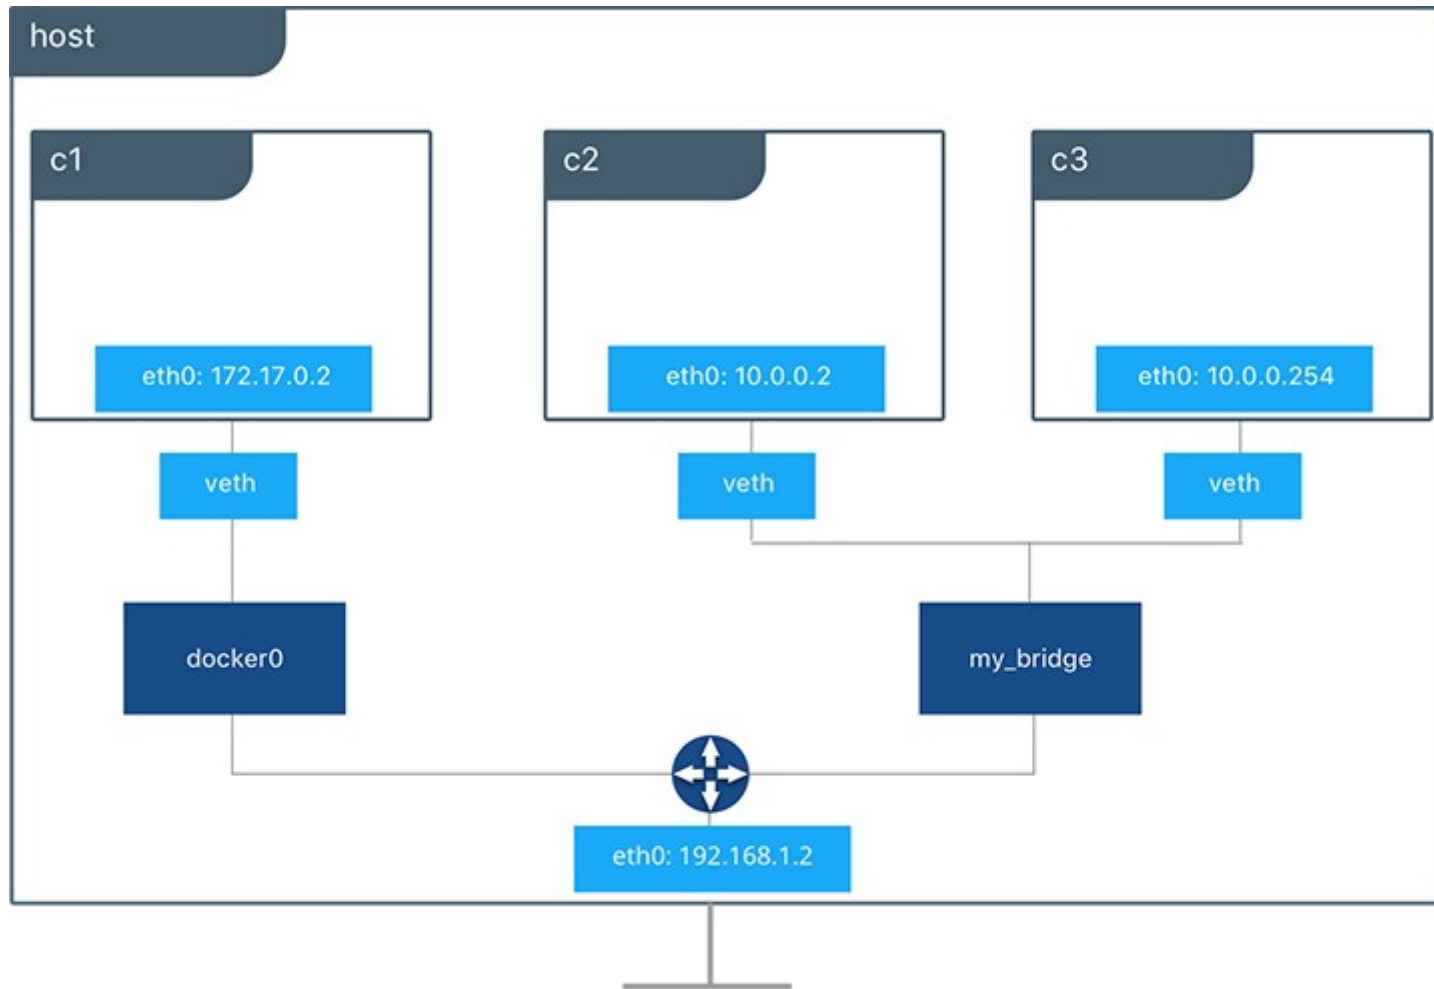
- **Docker provides multiple networking options**
  - Full overview: <https://docs.docker.com/network/>
  - All manipulate host's network configuration (by adding network devices, manipulating iptables) to:
    - create internal networks
    - connect containers to the internet
    - make containers to look like a physical device to the host OS
- **We will discuss only Bridge networks**

# Docker - networking

- **Bridge network**

- Created in terms of one local docker daemon (not for clusters)
- Containers within same bridge network can communicate with each other
- Containers in different networks cannot communication with each other
- Containers can call each other either with IPs or their names
- Default bridge network is created when Docker starts
  - All new containers are assigned to it unless specified otherwise
  - In the default bridge network containers may communicate only via IPs

# Docker - networking



# Docker - networking

- By default, all containers within a network expose all their ports to each other, but no ports to the outside world
- -> unless you specifically allow it, containers cannot be reached from the outside world
- Ports are exposed by binding HOST ports to container port

# Docker - networking

- Ports are exposed by binding HOST ports to container port
- **Example:**
  - Server hostname: myserver.com
  - Container running tomcat on port 8080
  - Port mapping 80 -> 8080
  - -> calls to <http://mysever.com/> will reach the Tomcat inside the container
  - -> calls to <http://myserver.com:8080/> will respond with „Server not available“

## Some more reading

<https://www.docker.com/resources/what-container>

<https://linuxcontainers.org/lxc>

<https://linuxcontainers.org/lxd/>

<https://pasztor.at/blog/lxc-vs-docker>

<https://goto.docker.com/rs/929-FJL-178/images/IDC-containerplatform-wp.pdf>

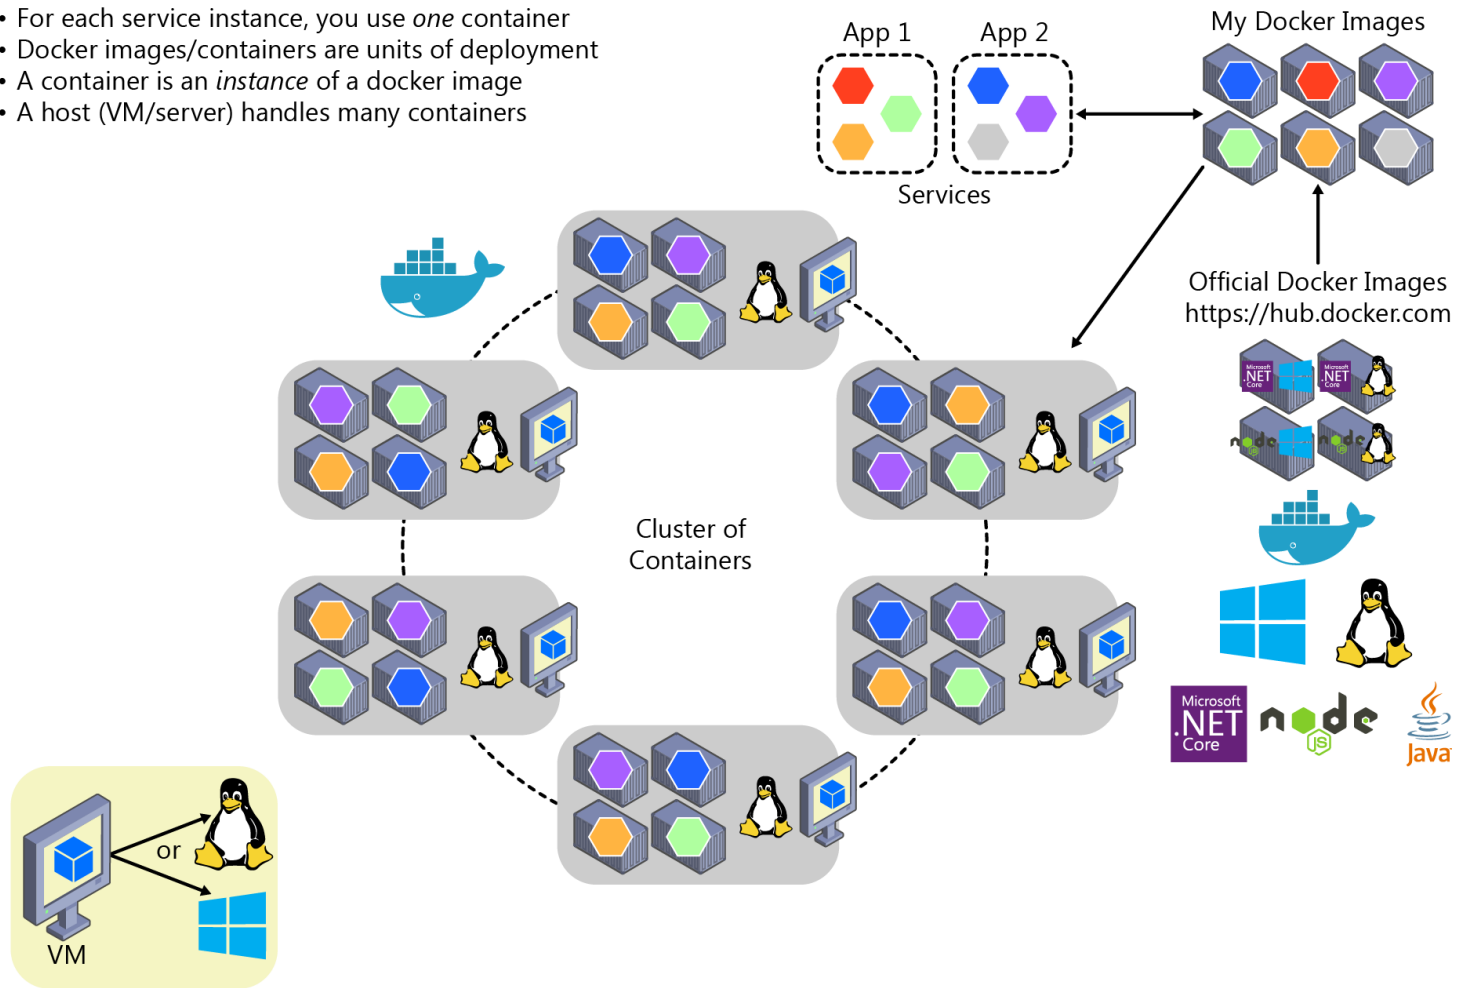
<https://success.docker.com/article/networking>

# Dockerizing Application

- **Application needs to be configurable via environment variables**
  - in theory it is possible to load configuration files into container via bind mounts, but that becomes somewhat complicated when running in cluster
- **That's it, you are good to go :)**
- **General rule is to have one process per docker container**
  - can be resolved by custom shell script or projects like <http://supervisord.org/>
  - but generally recommended approach is one process per container

# Dockerizing Application

- For each service instance, you use *one* container
- Docker images/containers are units of deployment
- A container is an *instance* of a docker image
- A host (VM/server) handles many containers



src: <https://docs.microsoft.com/en-us/dotnet/standard/containerized-lifecycle-architecture/design-develop-containerized-apps/orchestrate-high-scalability-availability>



# Container Orchestration

- **main strength of Docker ecosystem comes with easy creation of multi-service deployment and high-availability setup**
- **Many projects approach this**
  - Docker Compose
  - Kubernetes, Docker Swarm
  - AWS, Google Cloud, Azure...

# Docker Compose

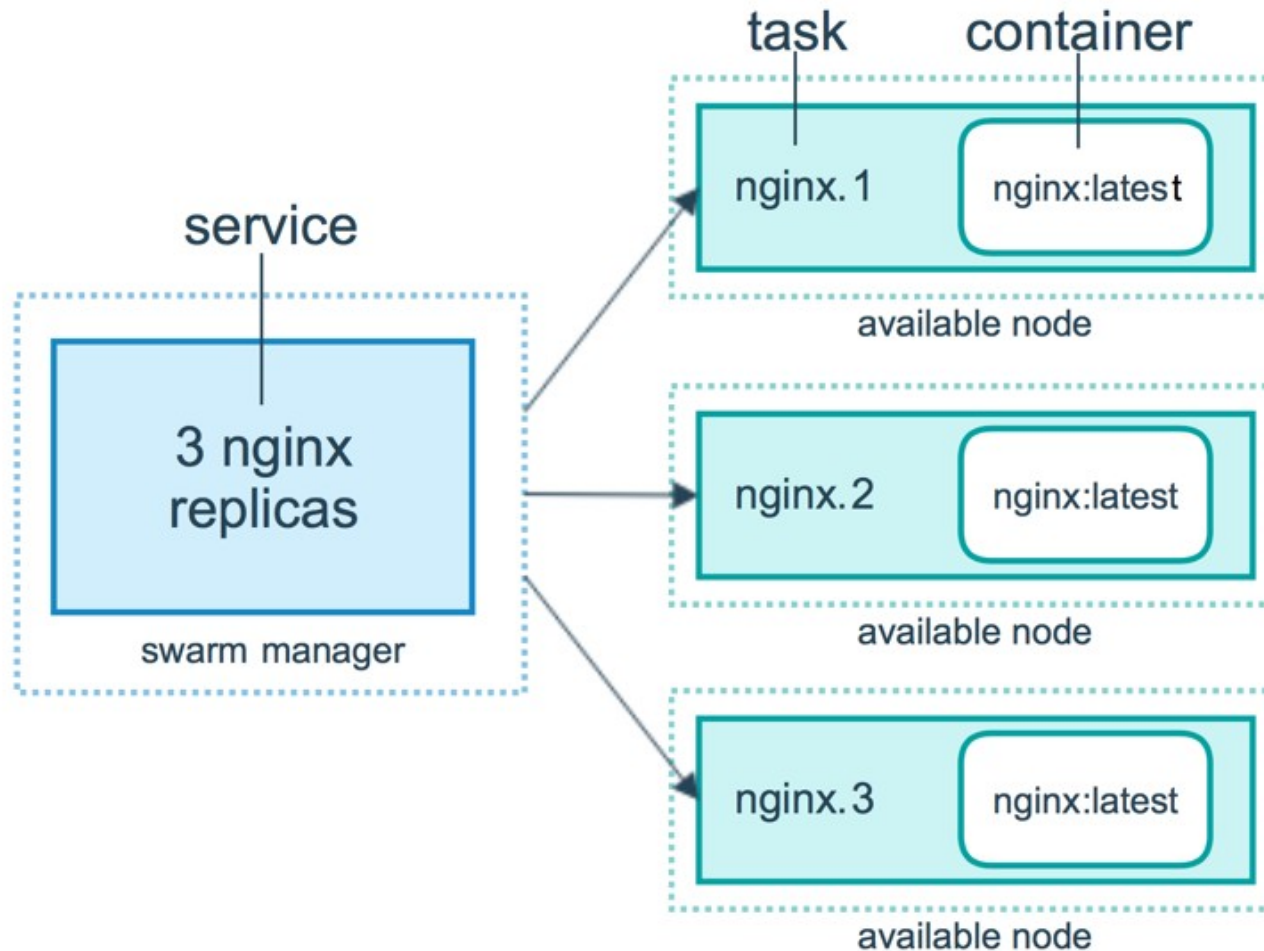
- **Tool for orchestrating containers based on single configuration file**
- **Easy to setup, easy to learn**
  - Can be used to configure services, networking, volumes, environment variables...
  - Good for development
  - Can be used for production on a small scale (single host)
    - No load-balancers
    - Needs OS tools to support restart in case of crash/server reboot
  - On larger scale, must be used together with Docker Swarm

# Docker Swarm

- **Cluster mode from Docker**

- Connects multiple Docker nodes into a cluster and allows their manipulation via single-point-of-entry API
- Developers use same commands (or docker-compose configuration) to deploy containers to the cluster
- Swarm handles deployment of individual container replicas to nodes

# Docker Swarm - Services



# Docker Swarm - Services

- Swarms runs internal DNS server
- Each service is given a record in DNS via which it can be reached
- Swarm handles load-balancing between the running replicas
- Configurable health-checks on service level
- This is a common approach provided by most of the orchestration tools
  - -> load-balancing out-of-the-box

# Kubernetes

- **More mature, based on Google's experience with containers**
  - also more widely used -> easier to find information
- **Provides extended features comparing to Docker Swarm**
  - auto-scaling, health-checks on container-level
  - is not restricted to Docker API (unlike Swarm which has to have the same API as Docker engine)

# Kubernetes vs Docker Swarm

- **Kubernetes cons**

- Do-it-yourself installation is rather complex
  - But there is a simple minikube tool for running on workstations and testing configurations
- Need to learn another tool

- **Docker Swarm pros**

- Simpler, same API as Docker

- **Still, Kubernetes is widely used and supported by main cloud providers**

# Openshift, GappEngine, AWS...

## Where can you host your Docker clusters?

- **Own VMs/servers**
  - -> awfull lot of work if you are serious about your data and availability
- **Openshift (IBM - RedHat)**
  - Cloud platform built on top of Kubernetes
  - Significant changes/improvements - support for builds, Jenkins integration, image registry and management
  - Very good visualisation
  - But: vendor lock-in



# Openshift, GappEngine, AWS...

## Where can you host your Docker clusters?

- **Google Cloud, AWS, Azure**
  - all provide options to run either plain Docker images with custom deployment configurations for the particular cloud
  - or they provide Kubernetes cluster hosting
- **Main con of all (including Openshift): price is much higher than with regular VMs**

# Final Thoughts

- **Note 1: The whole ecosystem evolves quickly, always check for the newest information**
  - Official documentation is your friend
- **Note 2: There is no such thing as silver bullet**
  - Not every application is suitable to run in container
  - But most of the time it is great at least for prototyping integration scenarios
- **Note 3: With the momentum given, containers seem to be the future, pay attention to them**
  - (this is author's personal opinion)

## Some more reading

<https://kubernetes.io/>

<https://cloud.google.com/kubernetes-engine/>

<https://aws.amazon.com/eks/>

<https://azure.microsoft.com/en-us/services/kubernetes-service/>



Questions?

