

Deployments, Operations

Jakub Danek
Yoso Czech s.r.o.

<mailto:jakub.danek@yoso.fi>
<https://www.linkedin.com/in/danekja/>

Motivation

- **Deployment significantly influences application ´s behaviour**
 - Security
 - Performance
 - Downtime during upgrades
- **Current trend is to minimize distance between development and operations → DevOps**

Topics

- What types of environments we have
- How can we deploy applications?
- Who is actually responsible for application deployments?
- Operations – backups, monitoring
- Examples – load-balanced and high-availability environments

Environments

- **Several different environments can be used during development cycle**
- **Each has different requirements on**
 - Hardware (performance)
 - Security
 - Stability
- **Each has different purpose**

Environments - Development

- **Used by developers to test the application outside of their workstation**
- **Performance requirements: low**
- **Stability requirements: low**
 - → developers can update the deployment several times a day, as much as needed
- **Security requirements: low**
 - No real data

Environments - Testing

- **Used by testers/product owners to validate/verify the implementation**
- **Performance requirements: low-medium**
 - Needs to respond in reasonable time, but very few users
- **Stability requirements: medium**
 - Testers/product owner need time to test new deployments, shouldn't change without notice
 - Sudden downtime does not cause any serious issues though
- **Security requirements: medium**
 - No real data, but setup should be production-like as much as possible

Environments - Pre-production

- **Used by testers/operations to test deployment in production-like environment, performance testing, security testing etc.**
- **Performance requirements: high**
 - Should copy production environment as much as possible
- **Stability requirements: med-high**
 - Tests deployment in production-like environment → deployments should happen based on predetermined schedule
 - Any issues causing downtime are a serious threat at this stage of development
- **Security requirements: high**
 - No real data, but same requirements as in production environment

Environments - Production

- **Used by customer**
- **Performance requirements: high**
 - Performance issues cause unhappy users
- **Stability requirements: high**
 - Deployments happen only at predetermined times
 - Users do not experience any downtime vs. down-time must be well communicated to users
- **Security requirements: high**
 - Real data, security vulnerabilities represent a serious threat

Environments - Other

- **Special testing environments can exist as well, e.g.:**
 - Performance testing
 - Maintaining production-like setup for testing can be expensive
 - → special cloud-based environment started only during the test-runs
 - Big feature test-env
 - Features taking long time to develop
 - Maybe large overhaul of the application is needed
 - → separate testing environment so that main development is not disrupted
 - → removed after the feature is finished and merged into main codebase

Naive deployment

- All done manually – build, deployment, configuration...
- No extra infrastructure needed
- Prone to human errors
- Slow

Scripted deployment

- **Improvement over naive approach – as much as possible automated with (shell) scripts**
 - Upload to server
 - Applying proper configuration
 - Restarting the service
 - Testing
- **But, developer is still required to exec those scripts on the server**

Automated deployment

- **All steps fully automated**
- **Designated build server**
 - Application responsible for all steps related to deployment
 - Deployment is triggered by clicking single button, after that no action is needed
 - Minimal human input → all deployments go through same process

Continuous Integration

- **Practice aiming at integrating new changes into main codebase as soon as possible**
- Limits overhead of integrating own changes and changes done by other team-members
- Requires high-quality testing process to ensure code stability
 - → build server can be configured to build and run all tests after every commit to master
 - → potential errors are discovered quickly
 - → developers can continue with their work while build server runs (timely) tests

Continuous Delivery

- **Practice aiming at deploying changes in small cycles**
 - A flexible approach to delivering new features/fixing critical errors
 - Possible due to small deployment overhead when using build server
- **Whole team needs to see whole build and deployment process (build, tests, reports) → easier to react**
- **Due to full automation, any version can be deployed to any environment at any time**
 - With constraints of course (e.g. you need to ensure no downtime of service during deployment to prod)

Build Servers

- **Jenkins**
 - Open-source, widely used, lots of plugins
- **TeamCity**
 - JetBrains
- **Travis CI**
 - CI server for GitHub hosted projects
- **Team Foundation Server**
 - Microsoft

Who is responsible?

- **Operations team (system administrators)**
 - Understand the servers, configuration
 - Lack detailed knowledge of the application
- **Developers**
 - Often do not want to deal with system administration
 - Have detailed knowledge of the application

DevOps

- **An approach/movement narrowing the gap between development and operations**
- **Developers and sysadmins need to cooperate**
 - Teams should be put together so that each team is able to
 - Deliver functionality
 - Test
 - Deploy
 - Maintain

- **Automation is crucial**
 - Limit manual work
 - Less error-prone
 - Fast delivery
- **Build and deployment automation**
 - We have just discussed that

- **Infrastructure automation**

- DevOps promotes same principles for infrastructure management as developers use
- Describe your infrastructure in configuration files
 - Can be versioned
 - Configuration management tools (Ansible, Puppet, Chef)

- **Virtualization**

- Cloud allows easy scaling and orchestration of new servers
- Virtualization tools/providers allow easy creation and management of identical server instances
- AWS, OpenStack, vagrant...

- **Containers**

- Another level of virtualization – allow shipping application together with its full runtime environment
- Build once, run anywhere
- Docker

DevOps

- <https://theagileadmin.com/what-is-devops/>
- <https://www.atlassian.com/devops>

Other Operations Tasks

- Security Configuration
- Monitoring
 - + scaling
- Backups

Security

- **It is common to deal with the following security tasks**
 - Firewall configurations
 - Certificate management
 - HTTPS configuration

Security

- **Firewall configuration**

- Allow connection only where needed, for example:

Application Server needs connection from internet on 80/443 ports (HTTP)

Database Server needs connection from the application server on 3306 port (MySQL)

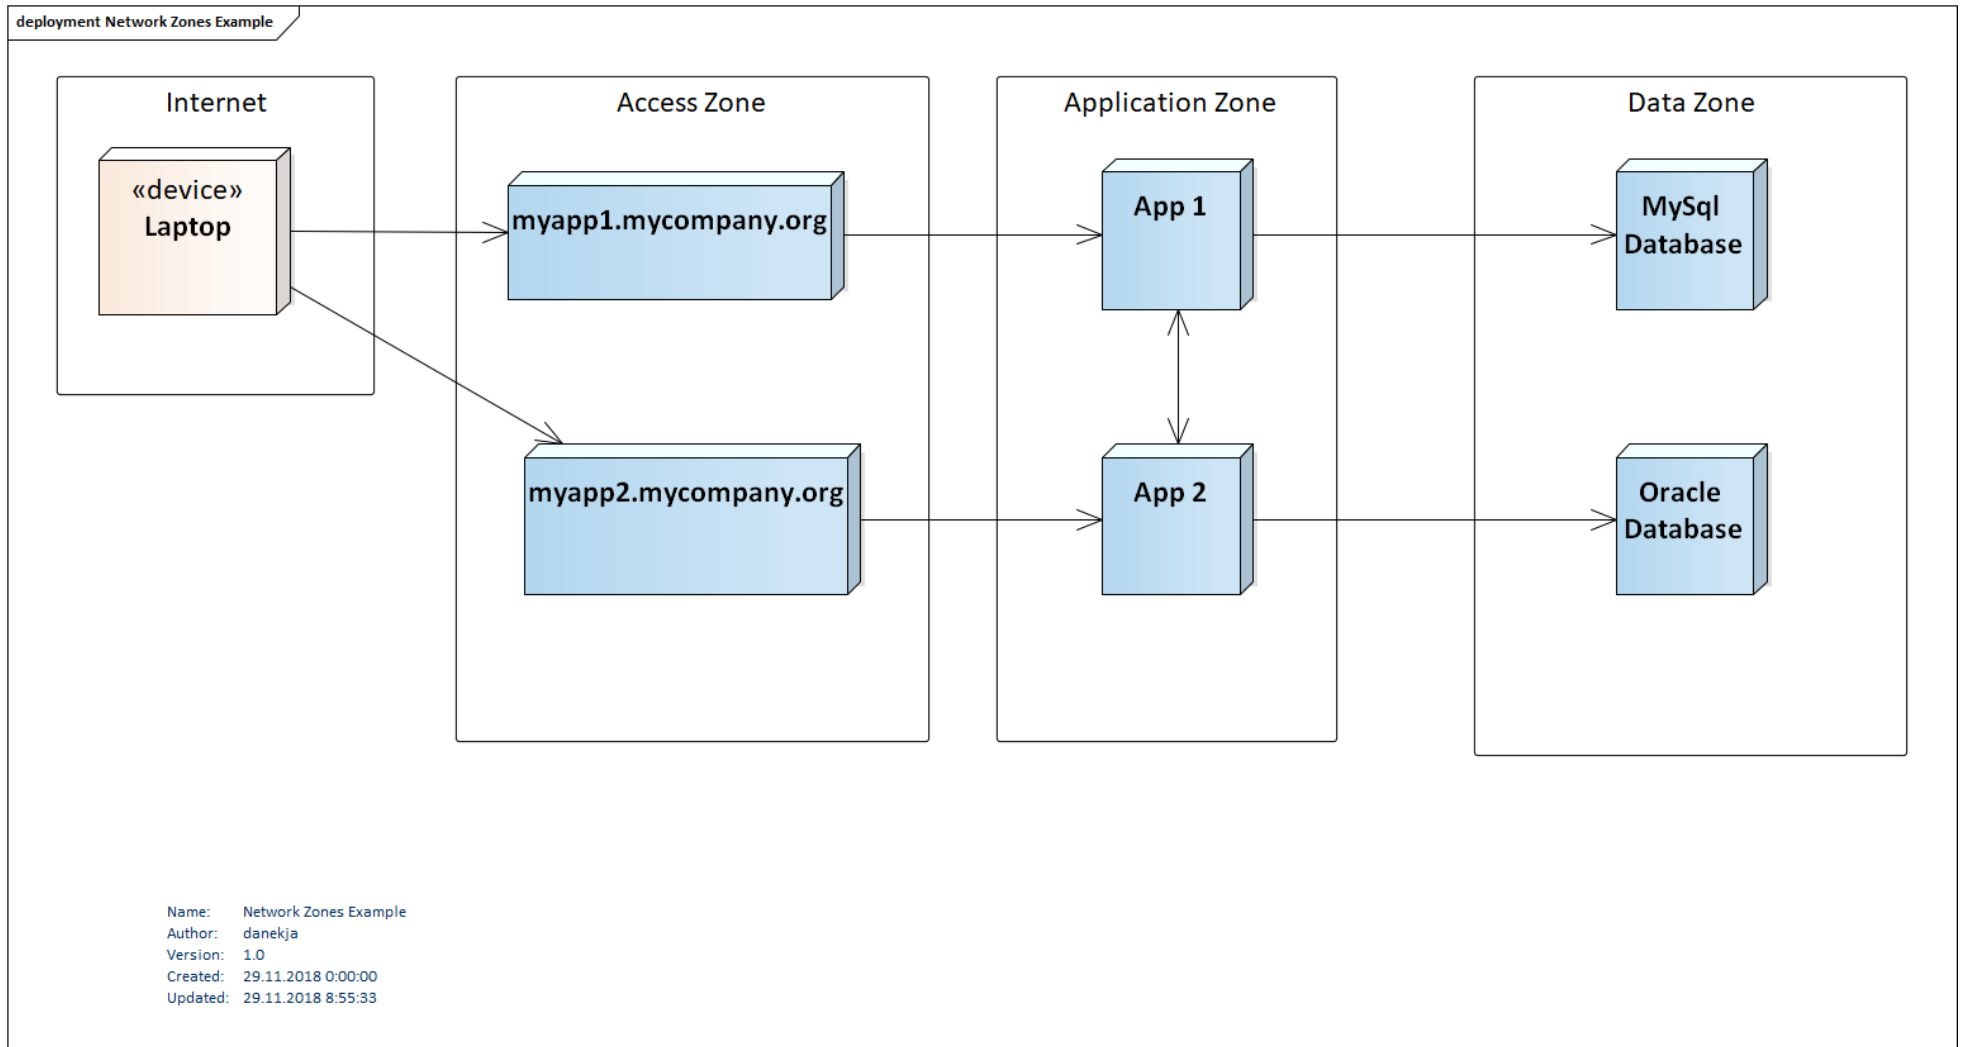
Both need connection on 22 port (SSH) from company's internal network.

Security

- **Firewall configuration**

- Companies with many servers tend to group them into network zones based on purpose
- → certain zones are unreachable from the internet
 - Implemented by firewall at entry point to the network
 - Useful for networks holding data stores (databases, file storage...)
- → it is necessary to configure proxies which have been granted access to communicate between zones
- → only a few points of access in space of hundreds of servers

Security – Network Zones Example



Security

Certificate Management

- Each certificate consists of public (certificate itself) and private (key) parts
- Public part has usually validity period (quite often one or two years)
 - After that it needs to be renewed

Security

Certificate Management – New Certificate

- **Private key is generated as random data**
 - Can be protected by passphrase, but not recommended for web certificates (you would need to solve the problem of providing the password each time web-server is restarted)
- **Using the private key, Certificate Signing Request (CSR) is made**
- **CSR is sent to certificate authority (CA) for signing**
 - CA verifies you own the domain for which you request certificate
- **Certificate file is received**

Certificate Management – Certificate Structure

- **Public Certificates defined in X.509 standard**
- **When creating CSR, several fields are asked:**
 - CN = Common Name = Fully Qualified Domain Name of your server
 - e.g. myapp1.mycompany.org
 - Browsers will reject the certificate if CN does not match the domain from which the certificate came
 - Country code, State, Organization, Organization Unit, email
 - Important for end-user to get information about your company, technically not necessary

Security

Server Configuration

- Web server needs access to the files with public and private parts

Example: nginx

```
server {  
    listen          443 ssl;  
    server_name     www.example.com;  
    ssl_certificate  /etc/pki/tls/certs/www.example.com.crt;  
    ssl_certificate_key /etc/pki/tls/private/www.example.com.key;  
    ssl_protocols   TLSv1 TLSv1.1 TLSv1.2;  
    ssl_ciphers     HIGH:!aNULL:!MD5;  
    ...  
}
```

Security

Questions?

Monitoring

- **You require information in order to react/prevent operation problems**
 - How much RAM has the server left?
 - How much free disk space there is?
 - Is the server under extensive load?
 - Does the web application answer to requests?

Monitoring

- **Checking manually when problems occur is too late!**
 - → monitoring application
- **External**
 - No special software required on the servers
 - Can monitor things like:
 - Application can be reached via network and answers requests
 - Response times of the application are within limits
 - Validity of HTTPS certificates

Monitoring






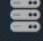











- **In-place**

- Monitoring software daemon installed on the server
- Can monitor things like:
 - CPU load
 - Memory use
 - Remaining free disk space

Monitoring

- All inputs from all servers aggregated at one place
- Dashboards with status
- Email/SMS/Chat messages if something goes wrong
 - If you realize your server is running out of disk space, you can react before users experience any issues
- Support for maintenance windows (no alarms are raised if you inform the monitor that downtime is due to planned maintenance)

Monitoring

 Site24x7 <input type="text" value="Search Monitors/Groups/Tags"/>		
<div>Monitor status  Last updated a few seconds ago</div>		
<div> Home</div>		
<div> Web</div>		
<div> APM</div>		
<div> Server</div>		
<div> VMware</div>		
<div> AWS</div>		
<div></div>		
Monitor Name ▾	Performance	Last Polled ▾
 brafitter-sizequery Lambda Function 	0 Count	4 minutes ago
 brafitter-update Lambda Function 	0 Count	2 minutes ago
 dynamodb SNSTopic 	-	5 minutes ago
 Membernet demo Website 	192 ms	2 minutes ago

Monitoring

<https://www.nagios.org/>

<https://icinga.com>

<https://www.site24x7.com/>

www.datadoghq.com

<https://prometheus.io/>

and many more...

Backups (and restores)

- **Data are the most valuable thing**
- **→ any production system needs to have backup mechanism**
- **In case of error data can be restored without too many losses**
 - Attack
 - Failed upgrade with broken data migration
 - SW/HW failure
 - User error

Backups (and restores)

- **Backup properties**

- Taking backup must not influence system runtime
- Can restore system to fully consistent state
- Stored at separate location

- **Full vs Incremental**

- Backup vs restore speed
- Space vs consistency threat





Examples



Examples – load-balanced environment

Nginx:

```
http {  
    upstream myapp {  
        server myapp1.intranet.mycompany.org;  
        server myapp2.intranet.mycompany.org;  
        server myapp3.intranet.mycompany.org;  
    }  
  
    server {  
        listen 443 ssl;  
        ...  
        location / {  
            proxy_pass http://myapp;  
        }  
    }  
}
```

Examples – load-balanced environment

HA Proxy:

frontend http

```
bind myapp.mycompany.org:443 ssl crt /etc/pki/tls/certs/myapp.pem
```

```
mode http
```

```
default_backend myapp
```

backend myapp

```
balance roundrobin
```

```
mode http
```

```
cookie MYAPPLBSRV insert indirect nocache maxidle 4h
```

```
server myapp1 myapp1.intranet.mycompany.org:80 check cookie myapp1
```

```
server myapp2 myapp2.intranet.mycompany.org:80 check cookie myapp2
```

```
server myapp3 myapp3.intranet.mycompany.org:80 check cookie myapp3
```

Examples – high-availability environment

- **Load-balancer solves two problems**

- Performance: Users can be distributed among multiple application nodes
- Stability: if one node dies, users automatically reach another, via same URL

- **What if load-balancer dies?**

- Virtual Router Redundancy Protocol (VRRP)



Examples – high-availability environment

- **Virtual Router Redundancy Protocol (VRRP)**
 - Creates Virtual Routers – group of multiple real routers
 - 1 active, rest backup
 - Routers within group communicate with each other
 - → if active router dies, new active is elected from the backups
- **How can we use this on application layer?**
 - Keepalived (<http://www.keepalived.org/>)

Examples – high-availability environment

- **Keepalived – need to configure the following on each node**
 - Virtual IP (address to which clients connect)
 - Always resolves to the active node
 - Virtual Router ID (unique VRRP group identifier within the same network segment, 0..255)
 - Check script/command
 - Keepalived calls this command regularly to check if the relevant service is running
 - In our example it would check e.g. if haproxy process is running
 - Priority: used during master elections
 - Password: used to secure communications between the instances